



Current font size: 10.95pt

0.1 Empirical Risk Minimization (ERM)

$$R(f) = \mathbb{E}_{(x,y) \sim \mathcal{D}} [\ell(f(x), y)] \quad \mathcal{D} \text{ data distrib.}$$

$$\text{Empirical Risk: } \hat{R}(f) = \frac{1}{n} \sum_{i=1}^n \ell(f(x_i), y_i)$$

0.2 Bias-Variance Tradeoff

$$\mathbb{E}[(y - \hat{y})^2] = \text{Bias}^2 + \text{Variance} + \text{Noise}$$

Variance: $\mathbb{E}[(\hat{y} - \mathbb{E}[\hat{y}])^2]$

0.3 Basic Loss Functions

$$\text{Logistic: } \ell(y, p) = -y \log p - (1-y) \log(1-p)$$

1 Representations

CRLB Lower bound on variance of unbiased estimators $\hat{\theta}$ of param θ . Assumes regularity: differentiable log-likelihood, finite variance.

Trick CRLB achieved iff estimator is efficient (e.g., MLE in exponential families) Multivariate: Use inverse Fisher matrix.

Hint for exams: Always check unbiasedness first; compute via Hessian or score function.

1.1 Formulas: Fisher Information

$$\text{Param sensitivity in likelihood. } I(\theta) = \mathbb{E} \left[\left(\frac{\partial}{\partial \theta} \log p(X|\theta) \right)^2 \right] = -\mathbb{E} \left[\frac{\partial^2 \log p}{\partial \theta^2} \log p(X|\theta) \right]$$

Multivariate ($\theta \in \mathbb{R}^k$): Matrix form

$$[I(\theta)]_{ij} = \mathbb{E} \left[\frac{\partial \log p}{\partial \theta_i} \frac{\partial \log p}{\partial \theta_j} \right] = -\mathbb{E} \left[\frac{\partial^2 \log p}{\partial \theta_i \partial \theta_j} \right]$$

Trick: For iid X_1, \dots, X_n , $I_n(\theta) = nI(\theta)$

Example: Gaussian $\mathcal{N}(\mu, \sigma^2 = 1)$

$$\text{Score: } \frac{\partial \log p}{\partial \mu} = x - \mu. \text{ Fisher: } I(\mu) = 1.$$

1.2 Rao-Cramér Lower Bound (CRLB)

$$\text{For unbiased } \hat{\theta}(X): \text{Var}(\hat{\theta}) \geq \frac{1}{I(\theta)} \quad (\text{scalar})$$

$$\text{Multiv. Var}(g(\hat{\theta})) \geq \left(\frac{\partial g}{\partial \theta} \right)^T I(\theta)^{-1} \left(\frac{\partial g}{\partial \theta} \right)$$

General CRLB: $\text{Cov}(\hat{\theta}) \succeq I(\theta)^{-1}$

Trick: Equality if $\hat{\theta} = a(\theta) \cdot s(X) + b(\theta)$, where $s(X)$ is sufficient statistic (e.g., in Gaussians, sample mean gives CRLB)

1.3 Calculus Recipes & Derivations

- Compute $I(\theta)$: (1) Write $\log L(\theta|X) = \sum \log p(x_i|\theta)$. (2) Take 2nd deriv or score sq. (3) Expectation over $p(X|\theta)$.
- For representations: Info in feature space: $I_\phi(\theta) = \mathbb{E}[\phi(X)^T \phi(X)]^{-1}$ (e.g., for linear models).
- Exam hint: CRLB bounds learning rates (e.g., variance in param est. for neural nets).

2 Gaussian Processes (GPs)

Definition Distribution over functions

$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$, where $m(\cdot)$ is mean, $k(\cdot, \cdot)$ is kernel (covariance function)

Key Kernels:

- RBF: $k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp \left(-\frac{\|\mathbf{x}-\mathbf{x}'\|^2}{2\ell^2} \right)$
- Linear: $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}' + c$.
- Matérn: $k(r) = \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\frac{\sqrt{2\nu}r}{\ell} \right)^\nu K_\nu \left(\frac{\sqrt{2\nu}r}{\ell} \right)$ ($\nu = 3/2$ or $5/2$ for smoothness).

GP Regression (Noisy Observations):

Train data $\mathbf{X}, \mathbf{y}, y = f(\mathbf{x}) + \epsilon, \epsilon \sim \mathcal{N}(0, \sigma_n^2)$

Prior: $\mathbf{f} \sim \mathcal{N}(\mathbf{0}, \mathbf{K})$, where $K_{ij} = k(x_i, x_j)$

Posterior predictive: $\mathbf{X}_*, \mathbf{f}_* | \mathbf{y} \sim \mathcal{N}(\bar{\mathbf{f}}_*, \text{cov}(\mathbf{f}_*))$

$$\begin{aligned} \bar{\mathbf{f}}_* &= \mathbf{K}_* \mathbf{x} (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}, \\ \text{cov}(\mathbf{f}_*) &= \mathbf{K}_{**} - \mathbf{K}_* \mathbf{x} (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{K}_{*}. \end{aligned}$$

$$\text{Marginal Likelihood } \log p(\mathbf{y}|\mathbf{X}, \theta) = -\frac{1}{2} \mathbf{y}^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y} - \frac{1}{2} \log |\mathbf{K} + \sigma_n^2 \mathbf{I}| - \frac{n}{2} \log 2\pi$$

3 Ensemble Methods

3.1 Bagging (Bootstrap Aggregating)

Average B models: $\hat{f}(\mathbf{x}) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(\mathbf{x})$

Reduction: $\text{Var}(\hat{f}) \approx \frac{1}{B} \text{Var}$ if uncorrelated

Random Forest: Bagging + random feature

Importance $I(f) = \sum_{\text{nodes}} \Delta \text{impurity} \cdot p(\text{node})$

3.2 Boosting

Sequential, weight misclassified points.

Final: $H(\mathbf{x}) = \text{sign}(\sum_m \alpha_m h_m(\mathbf{x}))$, $\alpha_m = \frac{1}{2} \log \frac{1-\epsilon_m}{\epsilon_m}$. Reduces bias.

AdaBoost

Weights $w_i^{(t+1)} = w_i^{(t)} \exp(-\alpha_t y_i h_t(\mathbf{x}_i))$, normalized. $\alpha_t = \frac{1}{2} \ln \left(\frac{1-\epsilon_t}{\epsilon_t} \right)$, where ϵ_t = weighted error.

Error bound: $\epsilon \leq 2^M \prod_m \sqrt{\epsilon_m(1-\epsilon_m)}$.

Gradient Boosting $\min L = \sum_i l(y_i, F(\mathbf{x}_i))$, update $F_m = F_{m-1} + \nu h_m$, where h_m fits pseudo-residuals $r_{im} = -\frac{\partial l}{\partial F_{m-1}(\mathbf{x}_i)}$.

4 Support Vector Machines (SVMs)

4.1 Hard-Margin SVM (Linearly Separable)

Primal: $\min \frac{1}{2} \|\mathbf{w}\|^2$ s.t. $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \forall i$

Dual: $\max \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$ s.t. $\alpha_i \geq 0, \sum_i \alpha_i y_i = 0$ Margin: $\gamma = \frac{2}{\|\mathbf{w}\|}$

Decision: $f(\mathbf{x}) = \text{sign}(\sum_i \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b)$

Support vectors: Points w. $\alpha_i > 0$ (on margin)

4.2 Soft-Margin SVM

Primal: $\min \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i \xi_i$ s.t. $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, \xi_i \geq 0$

Hinge loss: $\ell(y, \hat{y}) = \max(0, 1 - y\hat{y})$.

Dual: Same as hard-margin but $0 \leq \alpha_i \leq C$.

C trades bias/var. large $C \rightarrow$ hard-margin

4.3 Kernel Trick

Replace $\mathbf{x}_i^T \mathbf{x}_j$ with $k(\mathbf{x}_i, \mathbf{x}_j)$. Dual becomes: $\max \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)$

• Linear: $k(\mathbf{x}, \mathbf{z}) = \mathbf{x}^T \mathbf{z}$

• Polynomial: $k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z} + c)^d$

• RBF: $k(\mathbf{x}, \mathbf{z}) = \exp \left(-\frac{\|\mathbf{x}-\mathbf{z}\|^2}{2\sigma^2} \right)$

Mercer's condition: Kernel matrix ≥ 0 (PSD)

5 Neural Networks: Basics

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)}, \mathbf{a}^{(l)} = \sigma(\mathbf{z}^{(l)})$$

$$\text{Sigm. } \sigma(x) = \frac{1}{1+e^{-x}}, \sigma'(x) = \sigma(x)(1-\sigma(x))$$

$$\text{Cross-Entropy loss } \mathcal{L} = -\sum_i y_i \log(\hat{y}_i)$$

$$\text{Backpropagation: Output gradient: } \delta^{(L)} =$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(L)}} = (\hat{\mathbf{y}} - \mathbf{y}) \odot \sigma'(\mathbf{z}^{(L)})$$

$$\text{Hidden: } \delta^{(l)} = (\mathbf{W}^{(l+1)T} \delta^{(l+1)}) \odot \sigma'(\mathbf{z}^{(l)})$$

Weight update: $\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(l)}} = \delta^{(l)} \mathbf{a}^{(l-1)T}$. Trick: Use chain rule; initialize weights $\sim \mathcal{N}(0, \frac{2}{n_{in}})$

Optimization Tricks: Gradient Descent:

$$\theta \leftarrow \theta - \eta \nabla \mathcal{L}$$

Momentum: Add velocity term.

Adam: Adaptive learning rates with moments.

6 Attention Mechanisms

Scaled Dot-Product Attention

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left(\frac{\mathbf{Q} \mathbf{K}^T}{\sqrt{d_k}} \right) \mathbf{V}$$

\mathbf{Q} : Queries ($n \times d_k$), \mathbf{K} : Keys ($m \times d_k$), \mathbf{V} : Values ($m \times d_v$) - Trick: Scaling prevents softmax saturation; causal mask for decoders (upper triangle $-\infty$).

Multi-Head Attention

Concat($\text{head}_1, \dots, \text{head}_h$) \mathbf{W}^O , each head: $\text{head}_i = \text{Attention}(\mathbf{Q} \mathbf{W}_i^Q, \mathbf{K} \mathbf{W}_i^K, \mathbf{V} \mathbf{W}_i^V)$, $h = 8$ typical, allow parallel focus on subspaces

Self-Attention: $\mathbf{Q} = \mathbf{K} = \mathbf{V} = \mathbf{X} \mathbf{W}$

(input projection). Exam Tip: Captures dependencies without recurrence; $O(n^2)$ time.

7 Transformers

Architecture: Encoder (self-attn + FFN) stack; Decoder (masked self-attn + enc-dec attn + FFN) stack. - FFN: Two linear layers with ReLU: $\text{FFN}(\mathbf{x}) = \max(0, \mathbf{x} \mathbf{W}_1 + \mathbf{b}_1) \mathbf{W}_2 + \mathbf{b}_2$. - Residual: $\mathbf{x} \leftarrow \mathbf{x} + \text{Sublayer}(\mathbf{x})$. LayerNorm after.

Positional Encoding Add to input embeddings.

$$\text{PE}_{(pos, 2i+1)} = \sin \left| \cos \left(\frac{pos}{10000^{2i/d}} \right) \right|$$

Allows order awareness; fixed or learned.

8 Computer Vision

8.1 Convolutional Neural Networks (CNNs)

Key Concepts: Parameter sharing, local connectivity, translation invariance. Architectures: LeNet (simple), AlexNet (deep with ReLU/dropout).

Discrete Convolution (2D)

Input $I: (H, W, C)$, Kernel $K: (k_h \times k_w \times C)$

$$O[i, j] = \sum_{m=0}^{k_h-1} \sum_{n=0}^{k_w-1} \sum_{c=1}^C I[i+m, j+n, c] \cdot K[m, n, c] + b$$

Output size: $\lfloor (H - k_h + 2p)/s \rfloor + 1$, where p = padding, s = stride.

Pooling (Max/Avg): Reduces dims, e.g., max-pool: $O[i, j] = \max_{m,n} I[i \cdot s + m, j \cdot s + n]$.

Backpropagation in CNNs: Gradients via chain rule. For conv layer: - Weight grad:
 $\frac{\partial \mathcal{L}}{\partial K[m,n,c]} = \sum_{i,j} \frac{\partial \mathcal{L}}{\partial O[i,j]} \cdot I[i+m, j+n, c]$.
 - Input grad: Rotate kernel 180° and convolve with output grad.

9 Graph Neural Networks (GNNs)

9.1 Basics & Notation

Graph $G = (V, E)$, $|V| = n$ nodes, adjacency matrix $\mathbf{A} \in \{0, 1\}^{n \times n}$ (symmetric for undirected). Feature matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$ (node features). Degree matrix $\mathbf{D} = \text{diag}(\sum_j A_{ij})$.

Normalized adjacency: $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$ (self-loops), $\hat{\mathbf{A}} = \mathbf{D}^{-1/2} \tilde{\mathbf{A}} \mathbf{D}^{-1/2}$ (symmetric normalization).

Message passing: Update node v as $h_v^{(l+1)} = \sigma \left(\sum_{u \in \mathcal{N}(v)} m_{u \rightarrow v}^{(l)} \right)$, where m aggregates neighbor info.

9.2 Graph Convolutional Network (GCN)

Layer: $\mathbf{H}^{(l+1)} = \sigma(\hat{\mathbf{A}} \mathbf{H}^{(l)} \mathbf{W}^{(l)})$, with $\mathbf{H}^{(0)} = \mathbf{X}$.

Spectral view: Approximation of graph Laplacian $\mathbf{L} = \mathbf{D} - \mathbf{A}$, normalized $\hat{\mathbf{L}} = \mathbf{I} - \hat{\mathbf{A}}$.

9.3 Graph Attention Network (GAT)

Attention: $\alpha_{ij} = \text{softmax}_j(\text{LeakyReLU}(\mathbf{a}^\top [\mathbf{W}h_i \parallel \mathbf{W}h_j]))$

Update: $h_i^{(l+1)} = \sigma \left(\sum_{j \in \mathcal{N}(i) \cup i} \alpha_{ij} \mathbf{W}h_j^{(l)} \right)$.

Multi-head: Concat or average heads.

10 Information Theory

10.1 Key Measures

Entropy: $H(X) = -\mathbb{E}_{p(x)}[\log p(x)]$

Joint entropy: $H(X, Y) = -\mathbb{E}[\log p(x, y)]$.

Conditional: $H(Y|X) = H(X, Y) - H(X)$.

Mutual information: $I(X; Y) = H(X) + H(Y) - H(X, Y) = H(X) - H(X|Y) = \text{KL}(p(x, y) \| p(x)p(y)) \geq 0$.

Cross-entropy: $H(p, q) = -\mathbb{E}_p[\log q] = H(p) + \text{KL}(p\|q)$.

KL divergence $\text{KL}(p\|q) = \mathbb{E}_p[\log(p/q)] \geq 0$

Tricks: Jensen-Shannon divergence for stability: $\text{JSD}(p\|q) = \frac{1}{2} \text{KL}(p\|m) + \frac{1}{2} \text{KL}(q\|m)$, $m = (p+q)/2$.

11 Anomaly Detection

11.1 Statistical Methods

Z-Score: Score $z_i = \frac{x_i - \mu}{\sigma}$. Anomaly if $|z_i| > \theta$

Mahalanobis Distance: Accounts for cov.

$$D_M(\mathbf{x}) = \sqrt{(\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu})}$$

Anomaly if $D_M > \theta$ (e.g., from χ^2 dist.).

11.2 Proximity-Based Methods

11.3 Isolation Forest

Randomly partition until isolation. **Anomaly**

Score: $s(\mathbf{x}, n) = 2^{-\frac{E(h(\mathbf{x}))}{c(n)}}$, where $h(\mathbf{x})$ = path length, $E(\cdot)$ = avg over trees, $c(n) = 2H(n-1) - \frac{2(n-1)}{n}$ (H = harmonic number). Anomaly if $s \approx 0.5$ (normal) or $s \rightarrow 1$ (anomaly). Works well in high dims.

11.4 One-Class SVM

Hyperplane maximizing margin from origin

$$\min_{\mathbf{w}, \xi_i, \rho} \frac{1}{2} \|\mathbf{w}\|^2 + \frac{1}{\nu n} \sum_i \xi_i - \rho$$

$$\text{s.t. } \mathbf{w}^\top \phi(\mathbf{x}_i) \geq \rho - \xi_i$$

$$\text{Decision: } f(\mathbf{x}) = \text{sign}(\mathbf{w}^\top \phi(\mathbf{x}) - \rho). \nu \in (0, 1]$$

12 RL & Active Learning

12.1 Markov Decision Processes (MDPs)

$$V^\pi(s) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s \right]$$

$$\text{Action-value } Q^\pi(s, a) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a \right]$$

$$\text{Advantage } A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s).$$

$$\text{Bellman Expectation: } V^\pi(s) = \sum_a \pi(a|s) \sum_{s',r} P(s', r|s, a) [r + \gamma V^\pi(s')].$$

$$\text{Bellman Optimality: } V^*(s) = \max_a \sum_{s',r} P(s', r|s, a) [r + \gamma V^*(s')]$$

$$\text{Discounted Return: } G_t = \sum_{k=t}^{\infty} \gamma^{k-t} R_{k+1}.$$

$$\text{Policy Gradient Thm: } \nabla_\theta J(\theta) =$$

$$\mathbb{E}_\pi [\nabla_\theta \log \pi_\theta(a|s) Q^\pi(s, a)]$$

$$\text{REINFORCE: } \hat{\nabla} J = \sum_t \nabla_\theta \log \pi(a_t|s_t) G_t.$$

Var. reduct. Subtract baseline $b(s_t) \approx V(s_t)$

$$\text{Actor-Critic: } A = r + \gamma V(s') - V(s).$$

13 Counterfactual Invariance

SCM: $X_i = f_i(\mathbf{PA}_i, U_i)$, where \mathbf{PA}_i are parents, U_i noise. Intervention $\text{do}(X = x)$: Replace f_X with constant x . $P(Y|do(X = x)) = \sum_z P(Y|X = x, Z = z) P(Z|X = x)$
 Counterfactual: $P(Y_x = y|Y_{x'} = y')$

- Invariance: Model f is counterfactually invariant if $f(X, do(A)) = f(X)$ for action A (e.g., distribution shift robustness).

Invariance Condition

$$\mathbb{E}[Y|X, E] = \mathbb{E}[Y|X] \text{ for environment } E$$

14 Reproducing Kernel Hilbert Spaces (RKHS)

Definition 1 (RKHS). Hilbert space \mathcal{H} where eval $f \mapsto f(x)$ continuous. Reproducing: $f(x) = \langle f, K(\cdot, x) \rangle_{\mathcal{H}}$.

- Kernel ridge reg.: $\hat{f} = \underset{f}{\operatorname{argmin}} \|f\|_{\mathcal{H}}^2 + \frac{1}{n} \sum_i (y_i - f(x_i))^2$. Sol: $\alpha = (\mathbf{K} + \lambda \mathbf{I})^{-1} \mathbf{y}$.

15 Variational Autoencoders (VAEs)

15.1 ELBO Formula

$$\mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x) \| p(z))$$

16 Non-Parametric Bayesian Methods

16.1 Dirichlet Processes (DPs) & Infinite Mixtures

$G = \sum_{k=1}^{\infty} \pi_k \delta_{\theta_k}$, $\pi_k = v_k \prod_{j=1}^{k-1} (1 - v_j)$, $v_j \sim \beta(1, \alpha)$. Non-parametric alternative to finite GMMs; allows model complexity to grow with data.

17 PAC Learning

- Realizable** $\exists h^* \in \mathcal{H}$ with true risk $L(h^*) = 0$.

- PAC Learnable:** \exists learner s.t. \forall distributions \mathcal{D} , $\forall \epsilon, \delta > 0$, with prob. $\geq 1 - \delta$, outputs h with $L(h) \leq \epsilon$ using $m = m(\epsilon, \delta)$ samples.

- Agnostic PAC:** No assumption on h^* ; minimize excess risk over \mathcal{H} .

- True Risk:** $L(h) = \mathbb{E}_{(x,y) \sim \mathcal{D}} [\ell(h(x), y)]$ (e.g., 0-1 loss: $\ell = \mathbf{1}_{h(x) \neq y}$).

- Empirical Risk:** $\hat{L}(h) = \frac{1}{m} \sum_{i=1}^m \ell(h(x_i), y_i)$

17.1 VC Dimension & Shattering

- Shattering:** \mathcal{H} shatters set $S \subseteq \mathcal{X}$ if $|\{\mathbf{y} \in \{0, 1\}^{|S|} : \exists h \in \mathcal{H} \text{ realizes } \mathbf{y} \text{ on } S\}| = 2^{|S|}$.
- Growth Function:** $\Pi_{\mathcal{H}}(m) = \max_{S: |S|=m} |\{h_S : h \in \mathcal{H}\}| \leq \left(\frac{em}{d}\right)^d$ (Sauer-Shelah, if VC-dim $d < \infty$).

- VC Dimension** $d = \text{VC}(\mathcal{H})$: Largest $|S|$ s.t. \mathcal{H} shatters S (infinite if no such max).
- Trick for VC Calc:** Find largest shatterable set (e.g., for half-planes: 3 points not collinear shatter, 4 do not).

Fundamental Thm of PAC (Realizable, Finite \mathcal{H}): $m \geq \frac{1}{\epsilon} (\ln |\mathcal{H}| + \ln(1/\delta))$ samples suffice for $L(h) \leq \epsilon$ w.p. $\geq 1 - \delta$ via ERM.

Infinite \mathcal{H} (VC-based): For VC-dim d , $m \geq C \frac{d + \ln(1/\delta)}{\epsilon}$ (lower bound); upper: $m = O\left(\frac{d \ln(1/\epsilon) + \ln(1/\delta)}{\epsilon}\right)$.

Agnostic PAC (Uniform Convergence): w.p. $\geq 1 - \delta$, $|L(h) - \hat{L}(h)| \leq \sqrt{\frac{2d \ln(em/d) + \ln(2/\delta)}{m}}$

Sample Complexity (Agnostic): $m = O\left(\frac{d \ln(1/\epsilon) + \ln(1/\delta)}{\epsilon^2}\right)$ for excess risk $\leq \epsilon$.

Tricks:

- Use Hoeffding for finite $|\mathcal{H}|$: $\Pr(|L - \hat{L}| > \epsilon) \leq 2|\mathcal{H}|e^{-2m\epsilon^2}$.
- For VC, bound $\Pi_{\mathcal{H}}(m) \leq \sum_{i=0}^d \binom{m}{i} \leq (em/d)^d$.
- ERM is PAC if \mathcal{H} has finite VC-dim.