



Current font size: 10.95pt

1 Key ML Concepts

1.1 What is Machine Learning?

ML enables systems to learn from data without explicit programming. Goals: Prediction, inference, decision-making.

Types of ML:

- Supervised:** Labeled data; predict y from x (e.g., regression, classification).
- Unsupervised:** Unlabeled data; find patterns (e.g., clustering, dimensionality reduction).
- Reinforcement:** Agent learns via rewards (e.g., MDPs).
- Semi-supervised/Self-supervised:** Mix of labeled/unlabeled.

ML Pipeline: Data collection → Preprocessing → Feature engineering → Model selection → Training → Evaluation → Deployment.

1.2 Empirical Risk Minimization (ERM)

Core principle: Minimize average loss on training data as proxy for true risk.

Formulas:

- True Risk: $R(f) = \mathbb{E}_{(x,y) \sim \mathcal{D}}[\ell(f(x), y)]$, where ℓ is loss function, \mathcal{D} is data distribution.
- Empirical Risk: $\hat{R}(f) = \frac{1}{n} \sum_{i=1}^n \ell(f(x_i), y_i)$.
- Goal: $\hat{f} = \operatorname{argmin}_f \hat{R}(f)$ (often with regularization).

Trick: Overfitting occurs when $\hat{R} \ll R$; use validation sets, cross-validation.

1.3 Bias-Variance Tradeoff

Decomposes generalization error: $\mathbb{E}[(y - \hat{y})^2] = \text{Bias}^2 + \text{Variance} + \text{Noise}$.

Formulas:

- Bias: $\mathbb{E}[\hat{y}] - y$ (systematic error).
- Variance: $\mathbb{E}[(\hat{y} - \mathbb{E}[\hat{y}])^2]$ (sensitivity to data).

Exam Hint: High bias → underfitting (simple models); high variance → overfitting (complex models). Regularization reduces variance.

1.4 Basic Loss Functions

Common in early exams for supervised tasks.

- Regression (MSE): $\ell(y, \hat{y}) = (y - \hat{y})^2$.
- Classification (0-1 Loss): $\ell(y, \hat{y}) = \mathbb{I}(y \neq \hat{y})$.
- Logistic Loss: $\ell(y, p) = -y \log p - (1 - y) \log(1 - p)$ (for binary).

Trick: For optimization, use surrogates (e.g., hinge loss for SVM instead of 0-1).

1.5 Overfitting & Regularization

Prevent by adding penalty: $\hat{f} = \operatorname{argmin}_f \hat{R}(f) + \lambda \Omega(f)$, where Ω is complexity (e.g., L2: $\|w\|^2$).

Hint: Cross-validation for λ ; exams may ask to derive regularized ERM.

2 Representations

2.1 Key Concepts & Tricks

- Representations:** Feature maps $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$ for non-linear models (e.g., kernels). Exam trick: Use for bounds in estimation (e.g., info in embedded spaces).
- CRLB:** Lower bound on variance of unbiased estimators $\hat{\theta}$ of param θ . Assumes regularity: differentiable log-likelihood, finite variance.
- Trick: CRLB achieved iff estimator is efficient (e.g., MLE in exponential families). Multivariate: Use inverse Fisher matrix.
- Hint for exams: Always check unbiasedness first; compute via Hessian or score function.

2.2 Formulas: Fisher Information

Fisher info measures param sensitivity in likelihood. For scalar θ :

$$I(\theta) = \mathbb{E} \left[\left(\frac{\partial}{\partial \theta} \log p(X|\theta) \right)^2 \right] = -\mathbb{E} \left[\frac{\partial^2}{\partial \theta^2} \log p(X|\theta) \right]$$

Multivariate ($\theta \in \mathbb{R}^k$): Matrix form

$$[I(\theta)]_{ij} = \mathbb{E} \left[\frac{\partial \log p}{\partial \theta_i} \frac{\partial \log p}{\partial \theta_j} \right] = -\mathbb{E} \left[\frac{\partial^2 \log p}{\partial \theta_i \partial \theta_j} \right]$$

Trick: For iid samples X_1, \dots, X_n , $I_n(\theta) = nI(\theta)$.

Example: Gaussian $\mathcal{N}(\mu, \sigma^2 = 1)$

Score: $\frac{\partial \log p}{\partial \mu} = x - \mu$. Fisher: $I(\mu) = 1$.

2.3 Rao-Cramér Lower Bound (CRLB)

For unbiased estimator $\hat{\theta}(X)$:

$$\text{Var}(\hat{\theta}) \geq \frac{1}{I(\theta)} \quad (\text{scalar})$$

Multivariate (for function $g(\theta)$):

$$\text{Var}(g(\hat{\theta})) \geq \left(\frac{\partial g}{\partial \theta} \right)^T I(\theta)^{-1} \left(\frac{\partial g}{\partial \theta} \right)$$

General CRLB:

$$\text{Cov}(\hat{\theta}) \succeq I(\theta)^{-1}$$

Trick: Equality if $\hat{\theta} = a(\theta) \cdot s(X) + b(\theta)$, where $s(X)$ is sufficient statistic (e.g., in Gaussians, sample mean achieves CRLB).

2.4 Calculus Recipes & Derivations

- Compute $I(\theta)$: (1) Write $\log L(\theta|X) = \sum \log p(x_i|\theta)$. (2) Take 2nd deriv or score sq. (3) Expectation over $p(X|\theta)$.
- For representations: Info in feature space: $I_\phi(\theta) = \mathbb{E}[\phi(X)^T \phi(X)]^{-1}$ (e.g., for linear models).
- Exam hint: In ML, CRLB bounds learning rates (e.g., variance in param est. for neural nets).

Regularity Conditions: Support indep. of θ ; $\int p(x|\theta)dx = 1$ differentiable under integral.

3 Gaussian Processes (GPs)

Definition: GP is a distribution over functions $f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$, where $m(\cdot)$ is mean function (often 0), $k(\cdot, \cdot)$ is kernel (covariance function).

Key Kernels:

- RBF: $k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp\left(-\frac{\|\mathbf{x}-\mathbf{x}'\|^2}{2\ell^2}\right)$ (lengthscale ℓ , variance σ_f^2).

• Linear: $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}' + c$.	• For proofs: Derive exponential loss minimization for AdaBoost weights.	Trick: Boosting reduces bias; risk of overfitting—use early stopping.
• Matérn: $k(r) = \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\frac{\sqrt{2\nu}r}{\ell} \right)^\nu K_\nu \left(\frac{\sqrt{2\nu}r}{\ell} \right)$ ($\nu = 3/2$ or $5/2$ for smoothness).	5 Support Vector Machines (SVMs)	6.3 Key Tricks & Comparisons
GP Regression (Noisy Observations): $y = f(\mathbf{x}) + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma_n^2)$. Training data: \mathbf{X}, \mathbf{y} .	5.1 Hard-Margin SVM (Linearly Separable)	Bias-variance: Ensembles \otimes variance (bagging) or \otimes bias (boosting).
Prior: $\mathbf{f} \sim \mathcal{N}(\mathbf{0}, \mathbf{K})$, where $K_{ij} = k(x_i, x_j)$.	Primal: Minimize $\frac{1}{2} \ \mathbf{w}\ ^2$ s.t. $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \forall i$.	Error bound (AdaBoost): Training error $\leq \exp(-2 \sum_t (\frac{1}{2} - \epsilon_t)^2)$.
Posterior predictive: For test points \mathbf{X}_* , $\mathbf{f}_* \mathbf{y} \sim \mathcal{N}(\bar{\mathbf{f}}_*, \text{cov}(\mathbf{f}_*))$,	Dual: Maximize $\sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$ s.t. $\alpha_i \geq 0$, $\sum_i \alpha_i y_i = 0$.	Exam hint: For non-separable data, use soft-margin or kernels; ensembles for high-variance base learners (e.g., deep trees).
$\bar{\mathbf{f}}_* = \mathbf{K}_* \mathbf{x} (\mathbf{K}_{\mathbf{X}\mathbf{X}} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}$,	Decision: $f(\mathbf{x}) = \text{sign}(\sum_i \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b)$. Margin: $\gamma = \frac{2}{\ \mathbf{w}\ }$.	7 Neural Networks: Basics
$\text{cov}(\mathbf{f}_*) = \mathbf{K}_{**} - \mathbf{K}_* \mathbf{x} (\mathbf{K}_{\mathbf{X}\mathbf{X}} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{K}_{\mathbf{X}*}$.	Support vectors: Points where $\alpha_i > 0$ (on margin).	MLP Forward Pass: For layer l , $\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)}$, $\mathbf{a}^{(l)} = \sigma(\mathbf{z}^{(l)})$.
Marginal Likelihood (for hyperparams θ): $\log p(\mathbf{y} \mathbf{X}, \theta) = -\frac{1}{2} \mathbf{y}^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y} - \frac{1}{2} \log \mathbf{K} + \sigma_n^2 \mathbf{I} - \frac{n}{2} \log 2\pi$.	Primal: Minimize $\frac{1}{2} \ \mathbf{w}\ ^2 + C \sum_i \xi_i$ s.t. $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i$, $\xi_i \geq 0$.	Activation Functions: - Sigmoid: $\sigma(x) = \frac{1}{1+e^{-x}}$, $\sigma'(x) = \sigma(x)(1 - \sigma(x))$. - ReLU: $\sigma(x) = \max(0, x)$, derivative 1 if $x > 0$ else 0. Trick: Mitigates vanishing gradients. - Softmax: $\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$, for classification.
Tricks/Exam Hints:	Hinge loss: $\ell(y, \hat{y}) = \max(0, 1 - y\hat{y})$.	Loss Functions: - MSE: $\mathcal{L} = \frac{1}{2N} \sum_i (\hat{y}_i - y_i)^2$. - Cross-Entropy: $\mathcal{L} = -\sum_i y_i \log(\hat{y}_i)$. Hint: For multi-class, combine with softmax.
• Cholesky decomp for inversion: Stable for positive-definite \mathbf{K} .	Dual: Same as hard-margin but $0 \leq \alpha_i \leq C$.	Backpropagation: - Output gradient: $\delta^{(L)} = \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(L)}} = (\hat{\mathbf{y}} - \mathbf{y}) \odot \sigma'(\mathbf{z}^{(L)})$. - Hidden: $\delta^{(l)} = (\mathbf{W}^{(l+1)T} \delta^{(l+1)}) \odot \sigma'(\mathbf{z}^{(l)})$.
• Optimize θ via gradient descent on log-marginal (derivs w.r.t. ℓ, σ_f).	Trick: C trades bias/variance; large $C \otimes$ hard-margin.	- Weight update: $\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(l)}} = \delta^{(l)} \mathbf{a}^{(l-1)T}$. Trick: Use chain rule; initialize weights $\sim \mathcal{N}(0, \frac{2}{n_{in}})$ (He init for ReLU).
• GPs for classification: Use logistic/sigmoid + Laplace approx or EP.	5.3 Kernel Trick	Optimization Tricks: Gradient Descent: $\theta \leftarrow \theta - \eta \nabla \mathcal{L}$. Momentum: Add velocity term. Adam: Adaptive learning rates with moments.
• Scalability: Sparse GPs (e.g., FITC) approx large datasets.	Replace $\mathbf{x}_i^T \mathbf{x}_j$ with $k(\mathbf{x}_i, \mathbf{x}_j)$. Dual becomes: Maximize $\sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)$.	8 Attention Mechanisms
4 Ensembles	Common kernels:	Scaled Dot-Product Attention: $\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V}$. - \mathbf{Q} : Queries ($n \times d_k$), \mathbf{K} : Keys ($m \times d_k$), \mathbf{V} : Values ($m \times d_v$). - Trick: Scaling prevents softmax saturation; causal mask for decoders (upper triangle $-\infty$).
Key Methods:	• Linear: $k(\mathbf{x}, \mathbf{z}) = \mathbf{x}^T \mathbf{z}$	Multi-Head Attention: MultiHead = Concat(head ₁ , ..., head _h) \mathbf{W}^O , - Each head: head _i = $\text{Attention}(\mathbf{Q}\mathbf{W}_i^Q, \mathbf{K}\mathbf{W}_i^K, \mathbf{V}\mathbf{W}_i^V)$. - Hint: $h = 8$ typical; allows parallel focus on subspaces.
• Bagging (Bootstrap Aggregating): Train M models on bootstrap samples, average predictions. Reduces variance.	• Polynomial: $k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z} + c)^d$	Self-Attention: $\mathbf{Q} = \mathbf{K} = \mathbf{V} = \mathbf{X}\mathbf{W}$ (input projection). Exam Tip: Captures dependencies without recurrence; $O(n^2)$ time.
• Random Forest : Bagging + random feature subsets at splits. Importance: $I(f) = \sum_{\text{nodes}} \Delta \text{impurity} \cdot p(\text{node})$.	• RBF: $k(\mathbf{x}, \mathbf{z}) = \exp\left(-\frac{\ \mathbf{x}-\mathbf{z}\ ^2}{2\sigma^2}\right)$	9 Transformers
• Boosting (e.g., AdaBoost): Sequential, weight misclassified points. Final: $H(\mathbf{x}) = \text{sign}(\sum_m \alpha_m h_m(\mathbf{x}))$, $\alpha_m = \frac{1}{2} \log \frac{1-\epsilon_m}{\epsilon_m}$.	Merger's condition: Kernel matrix positive semi-definite.	Architecture: Encoder (self-attn + FFN) stack; Decoder (masked self-attn + enc-dec attn + FFN) stack. - FFN: Two linear layers with ReLU: $\text{FFN}(\mathbf{x}) = \max(0, \mathbf{x}\mathbf{W}_1 + \mathbf{b}_1)\mathbf{W}_2 + \mathbf{b}_2$.
• Gradient Boosting : Minimize loss $L = \sum_i l(y_i, F(\mathbf{x}_i))$, update $F_m = F_{m-1} + \nu h_m$, where h_m fits pseudo-residuals $r_{im} = -\frac{\partial l}{\partial F_{m-1}(\mathbf{x}_i)}$.	6 Ensemble Methods	
Bias-Variance Tradeoff: Ensembles reduce variance (bagging) or bias (boosting). Error bound for AdaBoost: $\epsilon \leq 2^M \prod_m \sqrt{\epsilon_m(1-\epsilon_m)}$.	6.1 Bagging (Bootstrap Aggregating)	
Tricks/Exam Hints:	Aggregate B bootstrapped models: $\hat{f}(\mathbf{x}) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(\mathbf{x})$ (regression) or majority vote (classification).	
• Diversity: Key to ensembles; measure via correlation of base learners.	Reduces variance: $\text{Var}(\hat{f}) \approx \frac{1}{B} \text{Var}(\text{single model})$ if uncorrelated.	
• Overfitting: Boosting can overfit; use early stopping or shrinkage $\nu < 1$.	Random Forest: Bagging + random feature subsets at splits. OOB error for validation.	
	6.2 Boosting	
	Sequential: Train weak learners on reweighted data.	
	AdaBoost: Weights $w_i^{(t+1)} = w_i^{(t)} \exp(-\alpha_t y_i h_t(\mathbf{x}_i))$, normalized. $\alpha_t = \frac{1}{2} \ln\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$, where ϵ_t = weighted error.	
	Final: $H(\mathbf{x}) = \text{sign}(\sum_t \alpha_t h_t(\mathbf{x}))$.	
	Gradient Boosting: Minimize loss by adding trees fitting residuals. Update: $f_m(\mathbf{x}) = f_{m-1}(\mathbf{x}) + \nu \cdot h_m(\mathbf{x})$ (ν : shrinkage).	

Residual: $\mathbf{x} \leftarrow \mathbf{x} + \text{Sublayer}(\mathbf{x})$. LayerNorm after.

Positional Encoding: $\text{PE}_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d}}\right)$, $\text{PE}_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d}}\right)$. - Added to input embeddings.

Trick: Allows order awareness; fixed or learned.

Training/Inference Tricks: Teacher forcing for training; beam search for generation. **Exam Focus:** Derive attention gradients or compare to RNNs (transformers handle long-range better).

10 Computer Vision

10.1 Convolutional Neural Networks (CNNs)

Key Concepts: Parameter sharing, local connectivity, translation invariance. Architectures: LeNet (simple), AlexNet (deep with ReLU/dropout).

Discrete Convolution (2D): For input I (size $H \times W \times C$), kernel K (size $k_h \times k_w \times C$),

$$O[i, j] = \sum_{m=0}^{k_h-1} \sum_{n=0}^{k_w-1} \sum_{c=1}^C I[i+m, j+n, c] \cdot K[m, n, c] + b$$

Output size: $\lfloor (H - k_h + 2p)/s \rfloor + 1$, where p = padding, s = stride.

Pooling (Max/Avg): Reduces dims, e.g., max-pool: $O[i, j] = \max_{m,n} I[i \cdot s + m, j \cdot s + n]$.

Activation Functions: ReLU: $f(x) = \max(0, x)$; Softmax for classification: $\sigma(z_i) = e^{z_i} / \sum e^{z_j}$.

10.2 Training & Optimization

Loss Functions: Cross-entropy for multi-class: $\mathcal{L} = -\sum y_i \log \hat{y}_i$.

Backpropagation in CNNs: Gradients via chain rule. For conv layer: - Weight grad: $\frac{\partial \mathcal{L}}{\partial K[m,n,c]} = \sum_{i,j} \frac{\partial \mathcal{L}}{\partial O[i,j]} \cdot I[i+m, j+n, c]$. - Input grad: Rotate kernel 180° and convolve with output grad.

Tricks: Batch norm: Normalize activations $\hat{x} = (x - \mu) / \sqrt{\sigma^2 + \epsilon}$, then $\gamma \hat{x} + \beta$. Dropout: Randomly zero neurons during training (prob p).

10.3 Applications & Metrics

Classification: Output via FC layer + softmax.

Object Detection Basics: Bounding boxes; IoU: $\text{IoU} = \frac{\text{Area}(A \cap B)}{\text{Area}(A \cup B)}$.

Segmentation: Pixel-wise classification; U-Net: Encoder-decoder with skip connections.

Exam Hints: Derive output shapes; explain why CNNs > FC nets (fewer params: $O(k^2C)$ vs. $O(HWC)$).

11 Graph Neural Networks (GNNs)

11.1 Basics & Notation

Graph $G = (V, E)$, $|V| = n$ nodes, adjacency matrix $\mathbf{A} \in \{0, 1\}^{n \times n}$ (symmetric for undirected). Feature matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$ (node features). Degree matrix $\mathbf{D} = \text{diag}(\sum_j A_{ij})$.

Normalized adjacency: $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$ (self-loops), $\hat{\mathbf{A}} = \mathbf{D}^{-1/2} \tilde{\mathbf{A}} \mathbf{D}^{-1/2}$ (symmetric normalization).

Message passing: Update node v as $h_v^{(l+1)} = \sigma\left(\sum_{u \in \mathcal{N}(v)} m_{u \rightarrow v}^{(l)}\right)$, where m aggregates neighbor info.

11.2 Graph Convolutional Network (GCN)

Layer: $\mathbf{H}^{(l+1)} = \sigma(\hat{\mathbf{A}} \mathbf{H}^{(l)} \mathbf{W}^{(l)})$, with $\mathbf{H}^{(0)} = \mathbf{X}$.

Spectral view: Approximation of graph Laplacian $\mathbf{L} = \mathbf{D} - \mathbf{A}$, normalized $\hat{\mathbf{L}} = \mathbf{I} - \hat{\mathbf{A}}$.

Over-smoothing: Deep GCNs make representations similar; mitigate with residual connections or normalization.

11.3 Graph Attention Network (GAT)

Attention: $\alpha_{ij} = \text{softmax}_j(\text{LeakyReLU}(\mathbf{a}^\top [\mathbf{W}h_i \| \mathbf{W}h_j]))$.

Update: $h_i^{(l+1)} = \sigma\left(\sum_{j \in \mathcal{N}(i) \cup i} \alpha_{ij} \mathbf{W}h_j^{(l)}\right)$. Multi-head: Concat or average heads.

11.4 Tasks & Pooling

Node classification: Predict labels from node embeddings.

Graph classification: Pool via readout $r = f(\{h_v | v \in V\})$ (e.g., mean, max, sum). Hierarchical pooling (e.g., DiffPool).

Link prediction: Score edges with $s(u, v) = h_u^\top h_v$ or MLP on concatenated embeddings.

Tricks: Use skip connections for deep GNNs; normalize features; handle heterophily with signed messages or higher-order neighbors.

12 Information Theory

12.1 Key Measures

Entropy: $H(X) = -\mathbb{E}_{p(x)}[\log p(x)] = -\sum p(x) \log p(x)$ (discrete); continuous: $-\int p(x) \log p(x) dx$.

Joint entropy: $H(X, Y) = -\mathbb{E}[\log p(x, y)]$.

Conditional: $H(Y|X) = H(X, Y) - H(X)$.

Mutual information: $I(X; Y) = H(X) + H(Y) - H(X, Y) = H(X) - H(X|Y) = \text{KL}(p(x, y) \| p(x)p(y)) \geq 0$.

Cross-entropy: $H(p, q) = -\mathbb{E}_p[\log q] = H(p) + \text{KL}(p \| q)$.

KL divergence: $\text{KL}(p \| q) = \mathbb{E}_p[\log(p/q)] \geq 0$, not symmetric.

12.2 Applications in ML/GNNs

Variational bound: ELBO in VAEs uses $\text{KL}(q \| p)$.

InfoMax in GNNs: Maximize $I(\mathbf{h}_v; \mathbf{h}_G)$ for unsupervised learning (e.g., InfoGraph).

Entropy regularization: In RL/policy, add $-H(\pi)$ to encourage exploration.

Chain rule: $H(X_1, \dots, X_n) = \sum H(X_i | X_{<i})$.

Tricks: Jensen-Shannon divergence for stability: $\text{JSD}(p \| q) = \frac{1}{2} \text{KL}(p \| m) + \frac{1}{2} \text{KL}(q \| m)$, $m = (p + q)/2$.

In GNNs: Use MI to measure information flow between layers or nodes.

13 Anomaly Detection

13.1 Key Concepts & Definitions

Anomaly (outlier): Data point deviating significantly from normal patterns. Types:

- **Unsupervised:** No labels; e.g., assume most data normal.
- **Semi-supervised:** Train on normal data only (one-class).
- **Supervised:** Labeled anomalies (rare due to imbalance).

Challenges: High dims (curse of dimensionality), imbalance, thresholding.

Tricks: Normalize data; use dimensionality reduction (e.g., PCA) pre-detection; evaluate with AUC-PR over AUC-ROC for imbalance.

13.2 Statistical Methods

Z-Score: Score $z_i = \frac{x_i - \mu}{\sigma}$. Anomaly if $|z_i| > \theta$ (e.g., 3).

Mahalanobis Distance: Accounts for covariance.

$$D_M(\mathbf{x}) = \sqrt{(\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu})}$$

Anomaly if $D_M > \theta$ (e.g., from χ^2 dist.).

13.3 Proximity-Based Methods

k-NN Outlier Score: Distance to k -th nearest neighbor $d_k(\mathbf{x})$.

Anomaly if $d_k > \theta$.

Local Outlier Factor (LOF): Compares local density.

1. Reachability dist.: $\text{rd}_k(p, o) = \max(d_k(o), d(p, o))$.
2. Local reach. density: $\text{lrd}_k(p) = \left(\frac{1}{N_k(p)} \sum_{o \in N_k(p)} \text{rd}_k(p, o)\right)^{-1}$.
3. LOF: $\text{LOF}_k(p) = \frac{1}{N_k(p)} \sum_{o \in N_k(p)} \frac{\text{lrd}_k(o)}{\text{lrd}_k(p)}$.

Anomaly if $\text{LOF} > 1$ (much lower local density).

13.4 Isolation Forest

Ensemble of isolation trees: Randomly partition until isolation.

Anomaly Score: $s(\mathbf{x}, n) = 2^{-\frac{E(h(\mathbf{x}))}{c(n)}}$, where $h(\mathbf{x})$ = path length, $E(\cdot)$ = avg over trees, $c(n) = 2H(n-1) - \frac{2(n-1)}{n}$ (H = harmonic number). Anomaly if $s \approx 0.5$ (normal) or $s \rightarrow 1$ (anomaly). Trick: Works well in high dims; depth limit for efficiency.

13.5 One-Class SVM

Hyperplane maximizing margin from origin (normal data).

$$\min_{\mathbf{w}, \xi_i, \rho} \frac{1}{2} \|\mathbf{w}\|^2 + \frac{1}{\nu n} \sum_i \xi_i - \rho$$

s.t. $\mathbf{w}^T \phi(\mathbf{x}_i) \geq \rho - \xi_i$. Decision: $f(\mathbf{x}) = \text{sign}(\mathbf{w}^T \phi(\mathbf{x}) - \rho)$. $\nu \in (0, 1]$ bounds outlier fraction.

13.6 Autoencoder-Based

Train on normal data; anomaly score = reconstruction error $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$. Threshold via validation set.

13.7 Evaluation

Anomaly Score Thresholding: Use quantiles or ROC curve. Metrics: Precision@K, AUC-PR (for imbalance).

14 RL & Active Learning

14.1 Markov Decision Processes (MDPs)

MDP: Tuple $(\mathcal{S}, \mathcal{A}, P, R, \gamma, \mu_0)$, where \mathcal{S} states, \mathcal{A} actions, $P(s'|s, a)$ transition prob., $R(s, a)$ reward, $\gamma \in [0, 1)$ discount, μ_0 initial state dist.

State-Value Fn: $V^\pi(s) = \mathbb{E}_\pi [\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) | s_0 = s]$.

Action-Value Fn: $Q^\pi(s, a) = \mathbb{E}_\pi [\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) | s_0 = s, a_0 = a]$.

Advantage Fn: $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$.

Bellman Expectation: $V^\pi(s) = \sum_a \pi(a|s) \sum_{s', r} P(s', r|s, a)[r + \gamma V^\pi(s')]$.

Bellman Optimality: $V^*(s) = \max_a \sum_{s', r} P(s', r|s, a)[r + \gamma V^*(s')]$. Trick: Optimal policy $\pi^*(s) = \arg\max_a Q^*(s, a)$.

Discounted Return: $G_t = \sum_{k=t}^{\infty} \gamma^{k-t} R_{k+1}$. Exam hint: Use for infinite-horizon problems; $\gamma < 1$ ensures convergence.

14.2 Value & Policy Iteration

Value Iteration (VI): Update $V(s) \leftarrow \max_a \mathbb{E}[R(s, a) + \gamma V(s')]$. Converges to V^* (contraction mapping, Banach fixed-point thm). Trick: Stop when $\max_s |V_{\text{new}}(s) - V_{\text{old}}(s)| < \epsilon(1-\gamma)/\gamma$.

Policy Iteration (PI): (1) Eval: Solve $V^\pi = (I - \gamma P^\pi)^{-1} R^\pi$. (2)

Improve: $\pi'(s) = \arg\max_a Q^\pi(s, a)$. Faster than VI for small \mathcal{A} .

Exam trick: PI is exact eval + greedy; VI is approx. eval + greedy. Derive from Bellman.

14.3 Model-Free RL (TD Methods)

TD(0) Update: $V(s) \leftarrow V(s) + \alpha[r + \gamma V(s') - V(s)]$ (TD error: $r + \gamma V(s') - V(s)$).

Q-Learning (Off-Policy): $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$. ϵ -greedy exploration.

SARSA (On-Policy): $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$.

Trick: Q-Learning converges to optimal even with suboptimal policy; SARSA to policy's Q. Exam: Compare bias/variance.

14.4 Policy Gradients & Actor-Critic

Policy Gradient Thm: $\nabla_\theta J(\theta) = \mathbb{E}_\pi[\nabla_\theta \log \pi_\theta(a|s) Q^\pi(s, a)]$.

Update: $\theta \leftarrow \theta + \alpha \nabla_\theta J$.

REINFORCE: Monte-Carlo est. $\hat{\nabla} J = \sum_t \nabla_\theta \log \pi(a_t|s_t) G_t$. Variance reduction: Subtract baseline $b(s_t) \approx V(s_t)$.

Actor-Critic: Actor updates policy; Critic est. V or Q via TD. A2C/A3C: Advantage $A = r + \gamma V(s') - V(s)$.

Trick: Softmax policy $\pi(a|s) = \exp(h_\theta(s, a)) / \sum_{a'} \exp(h_\theta(s, a'))$; gradient $\nabla \log \pi = \psi(s, a) - \sum_{a'} \pi(a'|s) \psi(s, a')$ (compat. features ψ).

Exam calc: Derive policy grad from $\frac{\partial}{\partial \theta} \mathbb{E}[R] = \mathbb{E}[R \nabla \log p(R|\theta)]$ (score fn trick).

14.5 Active Learning

Pool-based: Unlabeled pool \mathcal{U} , labeled \mathcal{L} . Query strategy selects $x^* \in \mathcal{U}$ to label.

Uncertainty Sampling: Query $x^* = \arg\max_x H(y|x, \mathcal{L})$ or $\arg\max_x [1 - P(\hat{y}|x)]$ (least confident). For binary: $\arg\max_x \min(P(y=1|x), P(y=0|x))$.

Query-by-Committee: Train committee of models; query max disagreement (vote entropy: $H = -\sum_y V(y)/C \log V(y)/C$, $V(y)$ votes for y).

Expected Model Change: Query $\arg\max_x \mathbb{E}_{y|x} [\|\nabla \ell(y, f(x))\|]$.

Trick: Balances exploration (uncertainty) vs. exploitation.

Exam: Compare to random sampling; derive entropy for multi-class.

Bayesian Active Learning: Use GP or BNN for $p(y|x)$; query max info gain $I(y; x) = H(y) - \mathbb{E}_{p(x)}[H(y|x)]$.

Exam hint: Active learning reduces labeling cost; focus on info-theoretic justifications.

15 Counterfactual Invariance

Key: Models robust to interventions (do-operator). Exam hint: Check invariance via causal graphs (SCMs); derive counterfactuals from joint distributions.

Definition 1 (Structural Causal Model (SCM)). SCM: $X_i = f_i(\mathbf{PA}_i, U_i)$, where \mathbf{PA}_i are parents, U_i noise. Intervention $\text{do}(X = x)$: Replace f_X with constant x .

Formulas:

$$P(Y|do(X = x)) = \sum_z P(Y|X = x, Z = z)P(Z|X = x)$$

Counterfactual : $P(Y_x = y|Y_{x'} = y')$ ("What if X was x given

Tricks/Hints:

- Invariance: Model f is counterfactually invariant if $f(X, do(A)) = f(X)$ for action A (e.g., distribution shift robustness).
- Exam proof: Use Pearl's ladder (observational \rightarrow interventional \rightarrow counterfactual). Check d-separation for identifiability.

Key Invariance Condition

Invariant if $\mathbb{E}[Y|X, E] = \mathbb{E}[Y|X]$ for environment E (no confounding).

16 Reproducing Kernel Hilbert Spaces (RKHS)

Key: Space \mathcal{H} of functions with kernel K . Exam hint: Prove properties via inner products; compute norms for regularization.

Definition 2 (RKHS). Hilbert space \mathcal{H} where eval $f \mapsto f(x)$ continuous. Reproducing: $f(x) = \langle f, K(\cdot, x) \rangle_{\mathcal{H}}$.

Formulas:

$$K(x, y) = \langle \phi(x), \phi(y) \rangle_{\mathcal{H}} \quad (\text{Kernel trick, feature map } \phi)$$

$$\langle K(\cdot, x), K(\cdot, y) \rangle_{\mathcal{H}} = K(x, y) \quad (\text{Reproducing property})$$

$$\|f\|_{\mathcal{H}}^2 = \sum_{i,j} \alpha_i \alpha_j K(x_i, x_j) \quad (\text{For } f(\cdot) = \sum_i \alpha_i K(\cdot, x_i))$$

$$\text{Mercer's Thm: } K(x, y) = \sum_{k=1}^{\infty} \lambda_k e_k(x) e_k(y) \quad (\text{Pos. def. kernel})$$

Tricks/Hints:

- Positive definite: K pos. def. if $\sum_{i,j} c_i c_j K(x_i, x_j) \geq 0 \forall \mathbf{c} \neq 0$.
- Kernel ridge reg.: $\hat{f} = \operatorname{argmin}_f \|f\|_{\mathcal{H}}^2 + \frac{1}{n} \sum_i (y_i - f(x_i))^2$.
Sol: $\alpha = (\mathbf{K} + \lambda \mathbf{I})^{-1} \mathbf{y}$.
- Exam: For new x_* , predict $f(x_*) = \mathbf{k}_*^T \alpha$, where $k_{*i} = K(x_*, x_i)$.

Common Kernels

Linear:	$K(x, y) = x^T y$
RBF:	$K(x, y) = \exp(-\ x - y\ ^2 / 2\sigma^2)$
Polynomial:	$K(x, y) = (x^T y + c)^d$

Methods

VAEs & Non-Parametric Bayesian

17 Variational Autoencoders (VAEs)

17.1 Key Concepts & Setup

VAE: Generative model with latent $z \sim \mathcal{N}(0, I)$. Encoder $q_\phi(z|x)$ (neural net) approximates posterior $p(z|x)$. Decoder $p_\theta(x|z)$ reconstructs x . Goal: Maximize marginal log-likelihood $\log p(x) = \mathbb{E}_{q(z|x)} [\log p(x|z)] - \text{KL}(q(z|x)\|p(z|x))$.

Trick: Use ELBO as surrogate (variational lower bound).

17.2 ELBO Formula

$$\text{ELBO}(\theta, \phi) = \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x)\|p(z))$$

- Reconstruction term: Measures data fit (e.g., Bernoulli or Gaussian likelihood). - KL term: Regularizes encoder to match prior (closed-form for Gaussians: $\text{KL}(\mathcal{N}(\mu, \sigma^2)\|\mathcal{N}(0, 1)) = \frac{1}{2} \sum (1 + \log \sigma^2 - \mu^2 - \sigma^2)$). - Exam hint: Derive by Jensen's inequality; $\log p(x) \geq \text{ELBO}$.

17.3 Reparameterization Trick

For backprop through sampling: $z = \mu + \sigma \odot \epsilon, \epsilon \sim \mathcal{N}(0, I)$. Allows gradient $\nabla_\phi \mathbb{E}_{q_\phi(z|x)} [f(z)] \approx \nabla_\phi f(\mu + \sigma \odot \epsilon)$.

Trick: Use for stochastic optimization; avoids high-variance Monte Carlo.

17.4 Training & Hints

- Loss: -ELBO (minimize via SGD/Adam). - β -VAE: Weight KL term by $\beta > 1$ for disentangled latents. - Exam-relevant: VAEs vs. GANs (VAEs stable, probabilistic); limitations (blurry outputs due to pixel-wise loss).

18 Non-Parametric Bayesian Methods

18.1 Gaussian Processes (GPs)

GP: $f(x) \sim \text{GP}(m(x), k(x, x'))$, mean $m(\cdot)$ (often 0), kernel $k(\cdot, \cdot)$ (e.g., RBF: $k(x, x') = \exp(-\|x - x'\|^2 / 2\ell^2)$).

Predictive Distribution (Regression, noisy $y = f(x) + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2)$): Let X_*, y_* for test, X, y for train. Kernel matrix $K = k(X, X) + \sigma^2 I$.

$$\mu_* = k(X_*, X)(K)^{-1}y, \quad \Sigma_* = k(X_*, X_*) - k(X_*, X)(K)^{-1}k(X, X_*)$$

- Posterior: $p(f_*|X_*, X, y) = \mathcal{N}(\mu_*, \Sigma_*)$. - Log-marginal likelihood: $\log p(y|X) = -\frac{1}{2}y^T K^{-1}y - \frac{1}{2} \log |K| - \frac{n}{2} \log 2\pi$ (for hyperparam tuning).

18.2 Kernels & Tricks

- Valid kernels: Positive semi-definite (e.g., linear, polynomial, Matérn). - Trick: Kernel trick for non-linear regression without explicit features. - Composition: Sum/product of kernels for flexibility. - Exam hint: Compute K matrix for small datasets; derive posterior mean/variance.

18.3 Dirichlet Processes (DPs) & Infinite Mixtures

DP: DP(α, H) for infinite mixture models (e.g., Dirichlet Process Mixture Model for clustering). Stick-breaking: $G = \sum_{k=1}^{\infty} \pi_k \delta_{\theta_k}, \pi_k = v_k \prod_{j=1}^{k-1} (1 - v_j), v_j \sim \beta(1, \alpha)$. - Exam-relevant: Non-parametric alternative to finite GMMs; allows model complexity to grow with data. - Hint: Chinese Restaurant Process analogy for sampling.

18.4 Hints for Non-Parametrics

- GPs vs. NNs: GPs exact uncertainty, but $O(n^3)$ cost (use sparse approximations). - Use for Bayesian optimization or regression with small data.

19 PAC Learning

19.1 Key Definitions & Framework

- Hypothesis class** \mathcal{H} : Set of functions $h : \mathcal{X} \rightarrow \mathcal{Y}$ (e.g., binary classifiers, $\mathcal{Y} = \{0, 1\}$).
- Realizable PAC**: $\exists h^* \in \mathcal{H}$ with true risk $L(h^*) = 0$.
- PAC Learnable**: \exists learner s.t. \forall distributions $\mathcal{D}, \forall \epsilon, \delta > 0$, with prob. $\geq 1 - \delta$, outputs h with $L(h) \leq \epsilon$ using $m = m(\epsilon, \delta)$ samples.
- Agnostic PAC**: No assumption on h^* ; minimize excess risk over \mathcal{H} .
- True/Generalization Risk**: $L(h) = \mathbb{E}_{(x,y) \sim \mathcal{D}} [\ell(h(x), y)]$ (e.g., 0-1 loss: $\ell = \mathbf{1}_{h(x) \neq y}$).
- Empirical Risk**: $\hat{L}(h) = \frac{1}{m} \sum_{i=1}^m \ell(h(x_i), y_i)$.

19.2 VC Dimension & Shattering

- Shattering**: \mathcal{H} shatters set $S \subseteq \mathcal{X}$ if $|\{\mathbf{y} \in \{0, 1\}^{|S|} : \exists h \in \mathcal{H} \text{ realizes } \mathbf{y} \text{ on } S\}| = 2^{|S|}$.
- Growth Function**: $\Pi_{\mathcal{H}}(m) = \max_{S:|S|=m} |\{h|_S : h \in \mathcal{H}\}| \leq \left(\frac{em}{d}\right)^d$ (Sauer-Shelah, if VC-dim $d < \infty$).
- VC Dimension** $d = \text{VC}(\mathcal{H})$: Largest $|S|$ s.t. \mathcal{H} shatters S (infinite if no such max).
- Examples**: VC(intervals on \mathbb{R}) = 2; VC(linear classifiers in \mathbb{R}^d) = $\lfloor d+1 \rfloor$; VC(axis-aligned rectangles in \mathbb{R}^2) = 4.
- Trick for VC Calc**: Find largest shatterable set (e.g., for half-planes: 3 points not collinear shatter, 4 do not).

19.3 Key Formulas & Bounds

Fundamental Thm of PAC (Realizable, Finite \mathcal{H}): $m \geq \frac{1}{\epsilon} (\ln |\mathcal{H}| + \ln(1/\delta))$ samples suffice for $L(h) \leq \epsilon$ w.p. $\geq 1 - \delta$ via ERM.

Infinite \mathcal{H} (VC-based): For VC-dim d , $m \geq C \frac{d + \ln(1/\delta)}{\epsilon}$ (lower bound); upper: $m = O\left(\frac{d \ln(1/\epsilon) + \ln(1/\delta)}{\epsilon}\right)$.

Agnostic PAC (Uniform Convergence): w.p. $\geq 1 - \delta$,

$$|L(h) - \hat{L}(h)| \leq \sqrt{\frac{2d \ln(em/d) + \ln(2/\delta)}{m}} \quad (\text{Rademacher/VC})$$

Sample Complexity (Agnostic): $m = O\left(\frac{d \ln(1/\epsilon) + \ln(1/\delta)}{\epsilon^2}\right)$ for excess risk $\leq \epsilon$.

Tricks:

- Use Hoeffding for finite $|\mathcal{H}|$: $\Pr(|L - \hat{L}| > \epsilon) \leq 2|\mathcal{H}| e^{-2m\epsilon^2}$.
- For VC, bound $\Pi_{\mathcal{H}}(m) \leq \sum_{i=0}^d \binom{m}{i} \leq (em/d)^d$.
- ERM is PAC if \mathcal{H} has finite VC-dim.